

Xue Bin (Jason) Peng
Year 4, PhD in Computer Science
jasonpeng142@hotmail.com
xbpeng.github.io

EDUCATION/AWARDS

- **PhD in Computer Science**, *University of California, Berkeley* 2017–Present
 - NSERC Postgraduate Scholarship
 - Berkeley Fellowship For Graduate Study
- **MSc in Computer Science**, *University of British Columbia* 2015 –2017
 - Governor-General’s Gold Medal
 - top of master’s class across all faculties (~2000 students)
 - NSERC Canada Graduate Scholarship Master’s Award
 - Theodore E Arnold Fellowship
 - CS Merit Award
- **Computer Science Honours**, *University of British Columbia* 2010 – 2015
 - Governor-General’s Silver Medal in Science
 - top of undergrad class in faculty of science (~2000 students)
 - Grade average: 95.6 %
 - Norman A M MacKenzie Scholarship
 - President’s Entrance Scholarship
 - Trek Excellence Scholarship
 - Greer Family Scholarship
 - Charles and Jane Banks Scholarship
 - Marie Kendall Memorial Scholarship in Science
 - Computer Science Scholarship
- **School of Interactive Arts and Technology**, *Simon Fraser University* 2009 – 2010
 - Gordon M. Shrum Scholarship

PUBLICATIONS

Refereed Journals/Conferences

- **Xue Bin Peng**, Erwin Coumans, Tingnan Zhang, Tsang-Wei Lee, Jie Tan, Sergey Levine. Learning Agile Robotic Locomotion Skills by Imitating Animals. *Robotics: Science and Systems (RSS)*, (2020). **Best paper**.
- Anirudh Goyal, Shagun Sodhani, Jonathan Binas, **Xue Bin Peng**, Sergey Levine, and Yoshua Benjio. Reinforcement Learning with Competitive Ensembles of Information-Constrained Primitives. *International Conference on Learning Representations (ICLR)*, (2020).
- Farzad Abdolhosseini, Hung Yu Ling, Zhaoming Xie, **Xue Bin Peng**, and Michiel van de Panne. On Learning Symmetric Locomotion. *Motion, Interaction and Games (MIG)*, (2019).
- **Xue Bin Peng**, Michael Chang, Grace Zhang, Pieter Abbeel, Sergey Levine. MCP: Learning Composable Hierarchical Control with Multiplicative Compositional Policies. *Neural Information Processing Systems (NeurIPS)*, (2019).
- **Xue Bin Peng**, Angjoo Kanazawa, Sam Toyer, Pieter Abbeel, and Sergey Levine. Variational Discriminator Bottleneck: Improving Imitation Learning, Inverse RL, and GANs by Constraining Information Flow. *International Conference on Learning Representations (ICLR)*, (2019).

- **Xue Bin Peng**, Angjoo Kanazawa, Jitendra Malik, Pieter Abbeel, and Sergey Levine. SFV: Reinforcement Learning of Physical Skills from Videos. *ACM Transactions on Graphics (Proc. SIGGRAPH Asia 2018)* 37, 6 (2018).
- **Xue Bin Peng**, Pieter Abbeel, Sergey Levine, and Michiel van de Panne. DeepMimic: Example-Guided Deep Reinforcement Learning of Physics-Based Character Skills. *ACM Transactions on Graphics (Proc. SIGGRAPH 2018)* 37, 4 (2018).
- **Xue Bin Peng**, Marcin Andrychowicz, Wojciech Zaremba, and Pieter Abbeel. Sim-to-Real Transfer of Robotic Control with Dynamics Randomization. *IEEE International Conference on Robotics and Automation (ICRA)*, (2018).
- **Xue Bin Peng**, Glen Berseth, KangKang Yin, and Michiel van de Panne. DeepLoco: Dynamic Locomotion Skills Using Hierarchical Deep Reinforcement Learning. *ACM Transactions on Graphics (Proc. SIGGRAPH 2017)* 36, 4 (2017).
- **Xue Bin Peng**, and Michiel van de Panne. Learning Locomotion Skills Using DeepRL: Does the Choice of Action Space Matter? *Proc. ACM SIGGRAPH / Eurographics Symposium on Computer Animation* (2017). **Best student paper.**
- **Xue Bin Peng**, Glen Berseth, and Michiel van de Panne. Terrain-adaptive locomotion skills using deep reinforcement learning. *ACM Transactions on Graphics (Proc. SIGGRAPH 2016)* 35, 4 (2016).
- **Xue Bin Peng**, Glen Berseth, and Michiel van de Panne. Dynamic Terrain Traversal Skills Using Reinforcement Learning. *ACM Transactions on Graphics (Proc. SIGGRAPH 2015)* 34, 4 (2015).

Non-Refereed

- Aviral Kumar, **Xue Bin Peng**, and Sergey Levine. Reward-Conditioned Policies. *arXiv preprint arXiv: 1912.13465* (2019).
- **Xue Bin Peng**, Aviral Kumar, Grace Zhang, and Sergey Levine. Advantage-Weighted Regression: Simple and Scalable Off-Policy Reinforcement Learning. *arXiv preprint arXiv: 1910.00177* (2019).

Posters and Abstracts

- **Xue Bin Peng**, Glen Berseth, and Michiel van de Panne. Learning Locomotion Skills Using DeepRL: Does the Choice of Action Space Matter? *NIPS Deep Reinforcement Learning Workshop*, (2016).
- **Xue Bin Peng**, Glen Berseth, and Michiel van de Panne. Terrain-adaptive locomotion skills using deep reinforcement learning. *NIPS Deep Learning Symposium*, (2016).
- **Xue Bin Peng**, Glen Berseth, and Michiel van de Panne. Dynamic Locomotion Across Variable Terrains Using Deep Reinforcement Learning. *Dynamic Walking*, (2016).
- **Xue Bin Peng**, Glen Berseth, and Michiel van de Panne. Dynamic Locomotion Skills for Obstacle Sequences Using Reinforcement Learning. *Dynamic Walking*, (2015).
- **Xue Bin Peng**, Glen Berseth, and Michiel van de Panne. Learning Dynamic Locomotion Skills for Terrains with Obstacles. *Reinforcement Learning and Decision Making*, (2015).

WORK EXPERIENCE

- Research Intern, Google Brain** June, 2019 – May, 2020
- Developed framework for learning locomotion skills from demonstrations for quadruped robots.
- Graduate Student Instructor, University of California, Berkeley** Jan – April 2019
- Taught tutorial sessions
 - Hosted office hours
 - Graded assignments and exams
- Member of Technical Staff (Intern), OpenAI** May – Aug, 2017
- Explored methods for transferring control policies from simulation to a physical robot
- Research Assistant, University of British Columbia** 2015 - 2017
- Developed methods to train motion control policies for physics-based character simulation
- Graduate Teaching Assistant, University of British Columbia** Jan – April 2017
- Hosted office hours
 - Wrote code for assignments
 - Graded assignments and exams
- Research Intern, Adobe Research** May – Aug, 2015
- Explored methods for physically-plausible motion control of simulated characters
- Lab Associate (Intern), Disney Research Pittsburgh** Jan – May, 2015
- Developed models of human gameplay strategies through imitation learning
 - Instrumented game to collect player data
- Undergraduate Teaching Assistant, University of British Columbia** 2011 - 2014
- Directed labs and hosted office hours
 - Wrote code for assignments
 - Graded assignments and exams
- Intern Software Developer, Microsoft Studios** May – Nov, 2013
- Developed real-time analytic approximation of area lights with different BRDFs
 - Implemented clustered forward lighting
 - Implemented environment map volumes and parallax correction
- Co-op Rendering Engineer, Capcom Vancouver** Jan – Aug, 2012
- Designed and created various rendering features through HLSL and C++
 - Designed a system for physically inspired image based lighting, utilizing real-time generation of dynamic environment maps
 - Implemented subsurface scattering for skin, distance field text and decal rendering, vertex animation, deferred lights, HDR cubemap support for Maya, and a variety of post-effects