

Policy Search

CMPT 729 G100


Jason Peng

Overview

- Policy Optimization
- Black Box Optimization
- Evolutionary Methods
- Finite-Difference Methods

Policy

$$\pi(\mathbf{a}|\mathbf{s})$$




Supervised Learning

$$\{(\mathbf{o}_0, \underline{\mathbf{a}}_0), (\mathbf{o}_1, \underline{\mathbf{a}}_1), \dots\}$$



Dataset




$$\min_{\pi} \mathbb{E}_{(\mathbf{o}, \mathbf{a}) \sim \mathcal{D}} [-\log \pi(\mathbf{a}|\mathbf{o})]$$

Behavioral Cloning

Supervised Learning


$$\{(\underline{\mathbf{x}}_i, \underline{y}_i)\}$$

Robot State




Motor Commands

?




Policy

$$\pi(\mathbf{a}|\mathbf{s})$$




Policy

$$\pi_{\underline{\theta}}(\mathbf{a}|\mathbf{s})$$



Optimization

$$\theta^* = \arg \max_{\theta} J(\pi_{\theta})$$




Just use gradient ascent!


Objective is often
NOT differentiable

~~$$\nabla_{\theta} J(\pi_{\theta})$$~~

Optimization




Optimization



$$\frac{\partial r_0}{\partial \theta}$$


Optimization



Reward function may not be differentiable


$$\frac{\partial r_0}{\partial \theta} = \overset{\text{X}}{\frac{\partial r_0}{\partial a_0}} \overset{\checkmark}{\frac{\partial a_0}{\partial \theta}}$$

Optimization



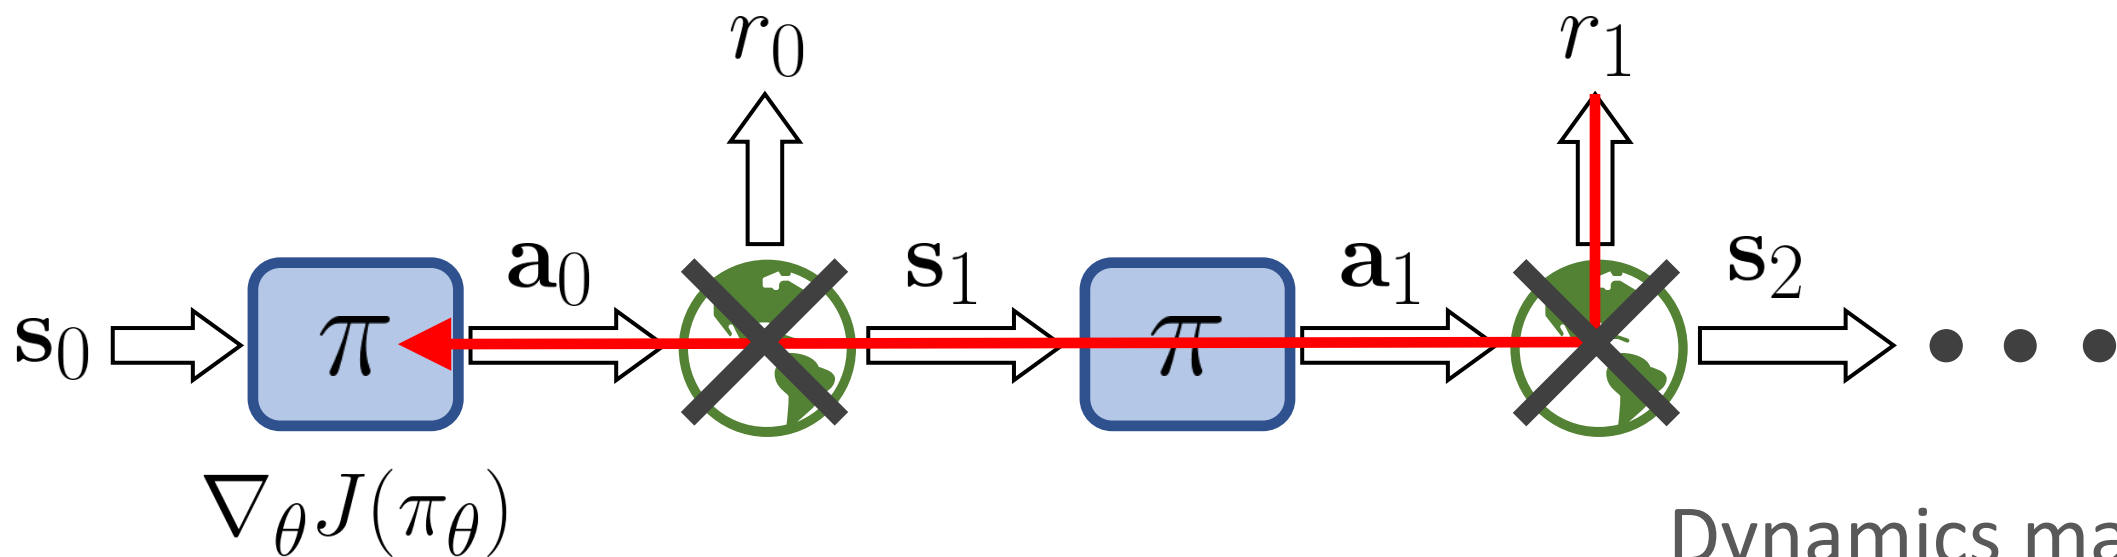
$$\frac{\partial r_1}{\partial \theta} = \overset{\text{X}}{\frac{\partial r_1}{\partial a_1}} \overset{\checkmark}{\frac{\partial a_1}{\partial \theta}}.$$

Optimization



$$\frac{\partial r_1}{\partial \theta} = \overset{\text{X}}{\frac{\partial r_1}{\partial \mathbf{a}_1}} \overset{\checkmark}{\frac{\partial \mathbf{a}_1}{\partial \theta}}.$$

Optimization




Dynamics may not be differentiable

$$\frac{\partial r_1}{\partial \theta} = \overset{\text{X}}{\frac{\partial r_1}{\partial \mathbf{a}_1}} \overset{\checkmark}{\frac{\partial \mathbf{a}_1}{\partial \theta}} + \overset{\text{X}}{\frac{\partial r_1}{\partial \mathbf{a}_1}} \overset{\checkmark}{\frac{\partial \mathbf{a}_1}{\partial s_1}} \boxed{\overset{\text{X}}{\frac{\partial s_1}{\partial \mathbf{a}_0}}} \overset{\checkmark}{\frac{\partial \mathbf{a}_0}{\partial \theta}} + \dots$$

Dynamics

Nondifferentiable Objective

$$\theta^* = \arg \max_{\theta} J(\pi_{\theta})$$



Just use gradient ascent!


Objective is often
NOT differentiable

~~$$\nabla_{\theta} J(\pi_{\theta})$$~~

Black Box Optimization


$$\theta^* = \arg \max_{\theta} \underline{J(\pi_{\theta})}$$

black box




Random Search

- Parameter space: $\theta \in \Theta$
- Each parameter is bounded: $\theta_i \in [\theta_i^{\min}, \theta_i^{\max}]$
- Idea: just sample randomly



Random Search

- Parameter space: $\theta \in \Theta$
- Each parameter is bounded: $\theta_i \in [\theta_i^{\min}, \theta_i^{\max}]$
- Idea: just sample randomly




Pick θ with highest expected return

Random Search for Hyper-Parameter Optimization
[Bergstra and Bengio 2012]

Grid Search

- Parameter space: $\theta \in \Theta$
- Each parameter is bounded: $\theta_i \in [\theta_i^{\min}, \theta_i^{\max}]$
- Idea: discretize space and evaluate



Pick θ with highest expected return

$$\theta^j \Rightarrow J(\pi_{\theta^j})$$

Random Search for Hyper-Parameter Optimization
[Bergstra and Bengio 2012]


Random Search

- ✓ Extremely simple
- ✓ Trivially parallel
- ✓ Can work well for small number of parameters (< 10)
- ✗ Curse of dimensionality
- ✗ Probably won't find an optimum

Smarter Search

- Adapt search samples base on objective


search distribution



sample

θ^j

evaluate



$J(\pi_{\theta^j})$


update



Smarter Search

- Adapt search samples base on objective


search distribution



sample

θ^j

evaluate




$J(\pi_{\theta^j})$

update

Random Search

- Random Search


search distribution



sample

θ^j

evaluate




$J(\pi_{\theta^j})$

update

Random Search

- Random Search


search distribution



sample

θ^j

evaluate



$J(\pi_{\theta^j})$


do nothing



Evolutionary Methods

- Adapt search samples base on objective


search distribution



sample

θ^j


evaluate



$J(\pi_{\theta^j})$


update

Cross-Entropy Method



The Cross-Entropy Method for Combinatorial and Continuous Optimization [Rubinstein 1999]
The Cross-Entropy Method for Optimization [Botev et al. 2013]

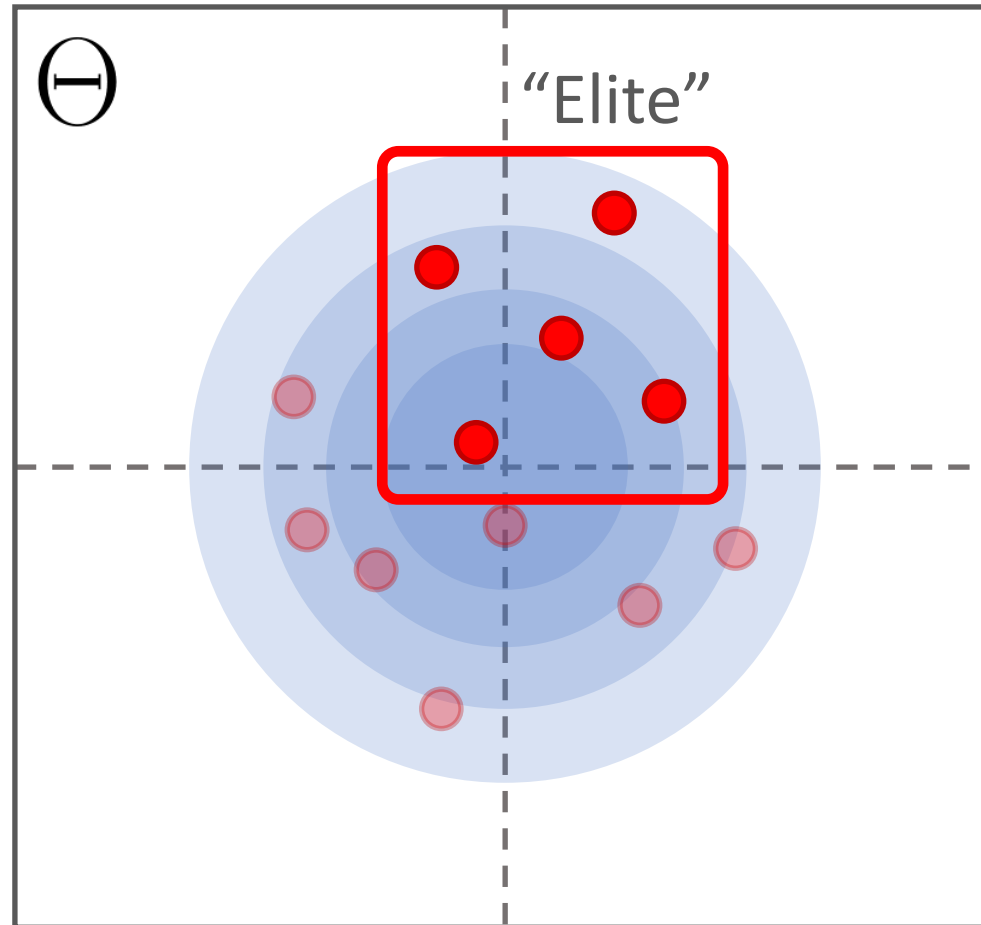
Cross-Entropy Method



The Cross-Entropy Method for Combinatorial and Continuous Optimization [Rubinstein 1999]

The Cross-Entropy Method for Optimization [Botev et al. 2013]


Cross-Entropy Method



The Cross-Entropy Method for Combinatorial and Continuous Optimization [Rubinstein 1999]

The Cross-Entropy Method for Optimization [Botev et al. 2013]


Cross-Entropy Method



The Cross-Entropy Method for Combinatorial and Continuous Optimization [Rubinstein 1999]

The Cross-Entropy Method for Optimization [Botev et al. 2013]


Cross-Entropy Method



The Cross-Entropy Method for Combinatorial and Continuous Optimization [Rubinstein 1999]

The Cross-Entropy Method for Optimization [Botev et al. 2013]


Cross-Entropy Method



The Cross-Entropy Method for Combinatorial and Continuous Optimization [Rubinstein 1999]

The Cross-Entropy Method for Optimization [Botev et al. 2013]

Cross-Entropy Method



The Cross-Entropy Method for Combinatorial and Continuous Optimization [Rubinstein 1999]

The Cross-Entropy Method for Optimization [Botev et al. 2013]

Cross-Entropy Method

ALGORITHM 4: CEM

- 1: $q^0 \leftarrow$ initialize search distribution
 - 2: **for** iteration $i = 0, \dots, k - 1$ **do**
 - 3: Sample parameters $\theta_1, \dots, \theta_n \sim q^i(\theta)$
 - 4: Evaluate performance of samples $J(\theta_1), \dots, J(\theta_n)$
 - 5: Select elite samples with highest performance $\hat{\theta}_1, \dots, \hat{\theta}_m$
 - 6: Update search distribution with elite samples:
 $q^{i+1} = \arg \max_q \frac{1}{m} \sum_{j=1}^m \log q(\hat{\theta}_j)$
 - 7: **end for**
 - 8: return best performing θ
-

Cross-Entropy Method

ALGORITHM 4: CEM

- 1: $q^0 \leftarrow$ initialize search distribution
 - 2: **for** iteration $i = 0, \dots, k - 1$ **do**
 - 3: Sample parameters $\theta_1, \dots, \theta_n \sim q^i(\theta)$
 - 4: Evaluate performance of samples $J(\theta_1), \dots, J(\theta_n)$
 - 5: Select elite samples with highest performance $\hat{\theta}_1, \dots, \hat{\theta}_m$
 - 6: Update search distribution with elite samples:
 $q^{i+1} = \arg \max_q \frac{1}{m} \sum_{j=1}^m \log q(\hat{\theta}_j)$
 - 7: **end for**
 - 8: return best performing θ
-

Cross-Entropy Method

ALGORITHM 4: CEM

- 1: $q^0 \leftarrow$ initialize search distribution
 - 2: **for** iteration $i = 0, \dots, k - 1$ **do**
 - 3: Sample parameters $\theta_1, \dots, \theta_n \sim q^i(\theta)$
 - 4: Evaluate performance of samples $J(\theta_1), \dots, J(\theta_n)$
 - 5: Select elite samples with highest performance $\hat{\theta}_1, \dots, \hat{\theta}_m$
 - 6: Update search distribution with elite samples:
 $q^{i+1} = \arg \max_q \frac{1}{m} \sum_{j=1}^m \log q(\hat{\theta}_j)$
 - 7: **end for**
 - 8: return best performing θ
-

Cross-Entropy Method

ALGORITHM 4: CEM

- 1: $q^0 \leftarrow$ initialize search distribution
 - 2: **for** iteration $i = 0, \dots, k - 1$ **do**
 - 3: Sample parameters $\theta_1, \dots, \theta_n \sim q^i(\theta)$
 - 4: Evaluate performance of samples $J(\theta_1), \dots, J(\theta_n)$
 - 5: Select elite samples with highest performance $\theta_1, \dots, \hat{\theta}_m$
 - 6: Update search distribution with elite samples:
 $q^{i+1} = \arg \max_q \frac{1}{m} \sum_{j=1}^m \log q(\hat{\theta}_j)$
 - 7: **end for**
 - 8: return best performing θ
-

Cross-Entropy Method

ALGORITHM 4: CEM

- 1: $q^0 \leftarrow$ initialize search distribution
 - 2: **for** iteration $i = 0, \dots, k - 1$ **do**
 - 3: Sample parameters $\theta_1, \dots, \theta_n \sim q^i(\theta)$
 - 4: Evaluate performance of samples $J(\theta_1), \dots, J(\theta_n)$
 - 5: Select elite samples with highest performance $\hat{\theta}_1, \dots, \hat{\theta}_m$
 - 6: Update search distribution with elite samples:
 $q^{i+1} = \arg \max_q \frac{1}{m} \sum_{j=1}^m \log q(\hat{\theta}_j)$
 - 7: **end for**
 - 8: return best performing θ
-

Cross-Entropy Method

ALGORITHM 4: CEM

- 1: $q^0 \leftarrow$ initialize search distribution
 - 2: **for** iteration $i = 0, \dots, k - 1$ **do**
 - 3: Sample parameters $\theta_1, \dots, \theta_n \sim q^i(\theta)$
 - 4: Evaluate performance of samples $J(\theta_1), \dots, J(\theta_n)$
 - 5: Select elite samples with highest performance $\hat{\theta}_1, \dots, \hat{\theta}_m$
 - 6: Update search distribution with elite samples:
 $q^{i+1} = \arg \max_q \frac{1}{m} \sum_{j=1}^m \log q(\hat{\theta}_j)$
 - 7: **end for**
 - 8: return best performing θ
-

Cross-Entropy Method

ALGORITHM 4: CEM

- 1: $q^0 \leftarrow$ initialize search distribution
 - 2: **for** iteration $i = 0, \dots, k - 1$ **do**
 - 3: Sample parameters $\theta_1, \dots, \theta_n \sim q^i(\theta)$
 - 4: Evaluate performance of samples $J(\theta_1), \dots, J(\theta_n)$
 - 5: Select elite samples with highest performance $\hat{\theta}_1, \dots, \hat{\theta}_m$
 - 6: Update search distribution with elite samples:
 $q^{i+1} = \arg \max_q \frac{1}{m} \sum_{j=1}^m \log q(\hat{\theta}_j)$
 - 7: **end for**
 - 8: return best performing θ
-

Cross-Entropy Method

ALGORITHM 4: CEM

- 1: $q^0 \leftarrow$ initialize search distribution
 - 2: **for** iteration $i = 0, \dots, k - 1$ **do**
 - 3: Sample parameters $\theta_1, \dots, \theta_n \sim q^i(\theta)$
 - 4: Evaluate performance of samples $J(\theta_1), \dots, J(\theta_n)$
 - 5: Select elite samples with highest performance $\hat{\theta}_1, \dots, \hat{\theta}_m$
 - 6: Update search distribution with elite samples:
 $q^{i+1} = \arg \max_q \frac{1}{m} \sum_{j=1}^m \log q(\hat{\theta}_j)$
 - 7: **end for**
 - 8: return best performing θ
-

$$q^i = \mathcal{N} \left(\mu^i, \Sigma^i \right)$$

$$\mu^i = \begin{pmatrix} \mu_1^i \\ \vdots \\ \mu_d^i \end{pmatrix}$$

$$\Sigma^i = \begin{bmatrix} \sigma_1^i & & \\ & \ddots & \\ & & \sigma_d^i \end{bmatrix}$$

Gaussian Maximum Likelihood

$$\arg \max_q \frac{1}{m} \sum_{j=1}^m \log q(\hat{\theta}_j) \quad \text{where } q = \mathcal{N}(\mu, \Sigma)$$

$$\nabla_q \frac{1}{m} \sum_{j=1}^m \log q(\hat{\theta}_j) = 0$$

$$\nabla_q \frac{1}{m} \sum_{j=1}^m -\frac{1}{2} (\hat{\theta}_j - \mu)^T \Sigma^{-1} (\hat{\theta}_j - \mu) - \frac{1}{2} \log \det(\Sigma) + C = 0$$

Gaussian Maximum Likelihood


$$\begin{aligned} \nabla_{\mu} \frac{1}{m} \sum_{j=1}^m -\frac{1}{2} \left(\hat{\theta}_j - \mu \right)^T \Sigma^{-1} \left(\hat{\theta}_j - \mu \right) - \frac{1}{2} \log \det (\Sigma) + C &= 0 \\ = \frac{1}{m} \sum_{j=1}^m \Sigma^{-1} \left(\hat{\theta}_j - \mu \right) &= 0 \\ = \frac{1}{m} \sum_{j=1}^m \underline{\Sigma} \Sigma^{-1} \left(\hat{\theta}_j - \mu \right) = \underline{\Sigma} 0 \end{aligned}$$

Gaussian Maximum Likelihood

$$\begin{aligned} & \nabla_{\mu} \frac{1}{m} \sum_{j=1}^m -\frac{1}{2} \left(\hat{\theta}_j - \mu \right)^T \Sigma^{-1} \left(\hat{\theta}_j - \mu \right) - \frac{1}{2} \log \det (\Sigma) + C = 0 \\ &= \frac{1}{m} \sum_{j=1}^m \Sigma^{-1} \left(\hat{\theta}_j - \mu \right) = 0 \\ &= \frac{1}{m} \sum_{j=1}^m \Sigma \Sigma^{-1} \left(\hat{\theta}_j - \mu \right) = \Sigma 0 \\ &= \frac{1}{m} \sum_{j=1}^m \left(\hat{\theta}_j - \mu \right) = 0 \\ &= \left(\frac{1}{m} \sum_{j=1}^m \left(\hat{\theta}_j \right) \right) - \mu = 0 \quad \Rightarrow \quad \boxed{\mu^* = \frac{1}{m} \sum_{j=1}^m \hat{\theta}_j} \end{aligned}$$

Gaussian Maximum Likelihood

$$\nabla_{\Sigma} \frac{1}{m} \sum_{j=1}^m -\frac{1}{2} \left(\hat{\theta}_j - \mu \right)^T \Sigma^{-1} \left(\hat{\theta}_j - \mu \right) - \frac{1}{2} \log \det(\Sigma) + C = 0$$



$$= \sum_{i=1}^d \log \sigma_i$$

$$\Sigma = \begin{bmatrix} \sigma_1 & & \\ & \ddots & \\ & & \sigma_d \end{bmatrix}$$

$$\nabla_{\sigma_i} \frac{1}{m} \sum_{j=1}^m -\frac{1}{2\sigma_i} \left(\hat{\theta}_{j,i} - \underline{\mu_i^*} \right)^2 - \frac{1}{2} \log \sigma_i = 0$$

Gaussian Maximum Likelihood

$$\nabla_{\sigma_i} \frac{1}{m} \sum_{j=1}^m -\frac{1}{2\sigma_i} \left(\hat{\theta}_{j,i} - \mu_i^* \right)^2 - \frac{1}{2} \log \sigma_i = 0$$

$$= \frac{1}{2\sigma_i^2} \frac{1}{m} \sum_{j=1}^m \left(\hat{\theta}_{j,i} - \mu_i^* \right)^2 - \frac{1}{2\sigma_i} = 0$$

$$\frac{1}{2\sigma_i^2} \frac{1}{m} \sum_{j=1}^m \left(\hat{\theta}_{j,i} - \mu_i^* \right)^2 = \frac{1}{2\sigma_i}$$

$$\sigma_i^* = \frac{1}{m} \sum_{j=1}^m \left(\hat{\theta}_{j,i} - \mu_i^* \right)^2$$

Cross-Entropy Method

ALGORITHM 4: CEM

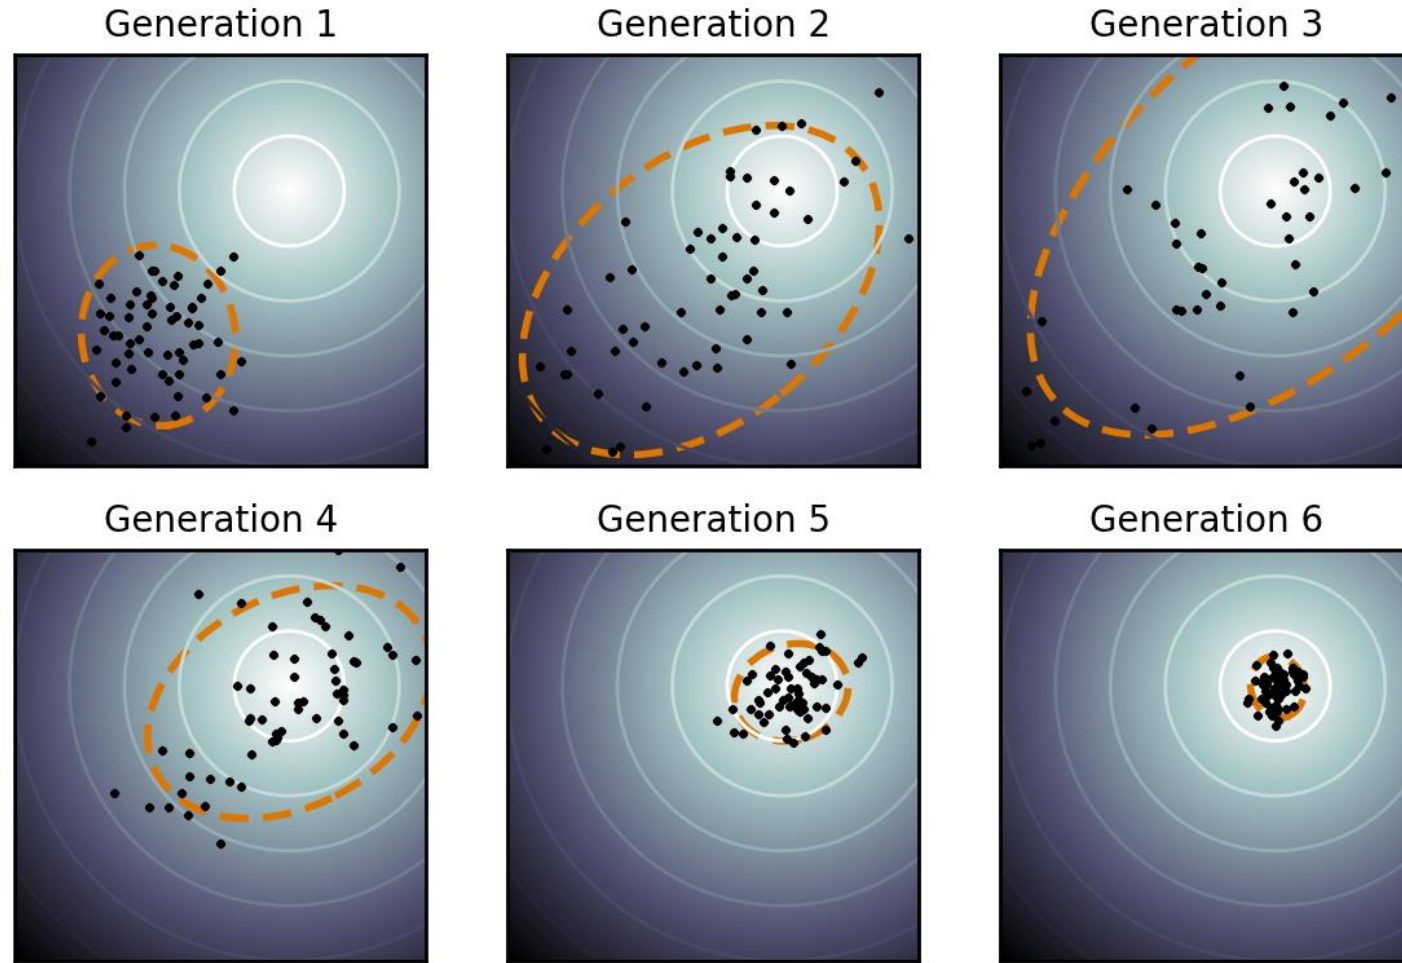
- 1: $q^0 \leftarrow$ initialize search distribution
 - 2: **for** iteration $i = 0, \dots, k - 1$ **do**
 - 3: Sample parameters $\theta_1, \dots, \theta_n \sim q^i(\theta)$
 - 4: Evaluate performance of samples $J(\theta_1), \dots, J(\theta_n)$
 - 5: Select elite samples with highest performance $\hat{\theta}_1, \dots, \hat{\theta}_m$
 - 6: Update search distribution with elite samples:
 $q^{i+1} = \arg \max_q \frac{1}{m} \sum_{j=1}^m \log q(\hat{\theta}_j)$
 - 7: **end for**
 - 8: return best performing θ
-

$$q^i = \mathcal{N} \left(\mu^i, \Sigma^i \right)$$

$$\mu^* = \frac{1}{m} \sum_{j=1}^m \hat{\theta}_j$$

$$\sigma_i^* = \frac{1}{m} \sum_{j=1}^m \left(\hat{\theta}_{j,i} - \mu_i^* \right)^2$$

Covariance Matrix Adaptation (CMA)




The CMA Evolution Strategy: A Comparing Review
[Hansen 2006]

Evolution Strategy Applications



Visual Foresight: Model-Based Deep Reinforcement Learning for Vision-Based Robotic Control
[Ebert et al. 2015]

Evolution Strategy Applications




Flexible Muscle-Based Locomotion for Bipedal Creatures
[Geijtenbeek et al. 2013]

Evolution Strategies

- ✓ Highly parallelizable
- ✓ Can work well for < 100 parameters
- ✗ Slow convergence
- ✗ Difficult to scale to large numbers of parameters

Nondifferentiable Objective

$$\theta^* = \arg \min_{\theta} J(\pi_{\theta})$$



$$J(\pi_{\theta})$$

nondifferentiable

$$\nabla_{\theta} J(\pi_{\theta})$$

Can we approximate?

Finite-Differences

- Start with initial guess θ^0
- Approximate partial derivatives using finite-differences

$$\frac{\partial J}{\partial \theta_j} \approx \frac{J(\pi_{\theta+\epsilon_j}) - J(\pi_{\theta-\epsilon_j})}{2\epsilon}$$


$$\epsilon_j = \begin{pmatrix} 0 \\ \vdots \\ \epsilon \\ \vdots \\ 0 \end{pmatrix}$$

jth component

very small


Finite-Differences

- Start with initial guess θ^0
- Approximate partial derivatives using finite-differences




Finite-Differences

- Start with initial guess θ^0
- Approximate partial derivatives using finite-differences




Finite-Differences


- Start with initial guess θ^0
- Approximate partial derivatives using finite-differences




Finite-Differences

- Start with initial guess θ^0

- Approximate partial derivatives using finite-differences
- 
- $$\Delta_j = \frac{J(\pi_{\theta+\epsilon_j}) - J(\pi_{\theta-\epsilon_j})}{2\epsilon}$$
- for every j

- Update: $\theta \leftarrow \theta + \alpha \Delta$
- 
- $$\Delta = \begin{pmatrix} \Delta_1 \\ \vdots \\ \Delta_n \end{pmatrix}$$

Finite-Differences



(a)



(b)



(c)



(d)



(e)



(f)

Policy Gradient Reinforcement Learning for Fast Quadrupedal Locomotion
[Kohl and Stone 2004]

Finite-Differences

- Start with initial guess θ^0

- Approximate partial derivatives using finite-differences

$$\Delta_j = \frac{J(\pi_{\theta+\epsilon_j}) - J(\pi_{\theta-\epsilon_j})}{2\epsilon}$$

- Update: $\theta \leftarrow \theta + \alpha \Delta$

for every j


2n evaluations
per iterations

Finite-Differences (Better)

- Start with initial guess θ^0
- Sample direction vector δ
- Approximate directional derivative


$$\Delta = \frac{J(\pi_{\theta+\epsilon\delta}) - J(\pi_{\theta-\epsilon\delta})}{2\epsilon} \delta$$

} fewer evaluations
per iteration

- 
- Update: $\theta \leftarrow \theta + \alpha\Delta$

— “directional derivative”



Augmented Random Search (ARS)



| Task | # timesteps | Maximum average reward after # timesteps | | | | |
|----------------|-------------|--|----------------|----------------|----------------|----------------|
| | | ARS | PPO | A2C | CEM | TRPO |
| Swimmer-v1 | 10^6 | 361 | ≈ 110 | ≈ 30 | ≈ 0 | ≈ 120 |
| Hopper-v1 | 10^6 | 3047 | ≈ 2300 | ≈ 900 | ≈ 500 | ≈ 2000 |
| HalfCheetah-v1 | 10^6 | 2345 | ≈ 1900 | ≈ 1000 | ≈ -400 | ≈ 0 |
| Walker2d-v1 | 10^6 | 894 | ≈ 3500 | ≈ 900 | ≈ 800 | ≈ 1000 |

Simple Random Search Provides a Competitive
Approach to Reinforcement Learning
[Mania et al. 2018]

ARS Applications



Policies Modulating Trajectory Generators
[Isken et al. 2018]

Summary

- Policy Optimization
- Black Box Optimization
- Evolutionary Methods
- Finite-Difference Methods