

# **Model-Based Reinforcement Learning**

CMPT 729 G100

Jason Peng

# Overview

---

- Model-Based RL
- DYNA
- Model Representations
- Uncertainty Estimation
- MPC


# Taxonomy of RL Algorithms

---

- Policy-Based Methods
- Value-Based Methods
- Actor-Critic Methods
- **Model-Based Methods**


# Reinforcement Learning

---



# RL Algorithms

---



# Sample Complexity

---



Simulation


Learning Agile Robotic Locomotion Skills by Imitating Animals  
[Peng et al. 2020]

# Sample Complexity

---



Simulation



Learning Agile Robotic Locomotion Skills by Imitating Animals  
[Peng et al. 2020]

# Sample Complexity



Simulation




Real World

Learning Agile Robotic Locomotion Skills by Imitating Animals  
[Peng et al. 2020]



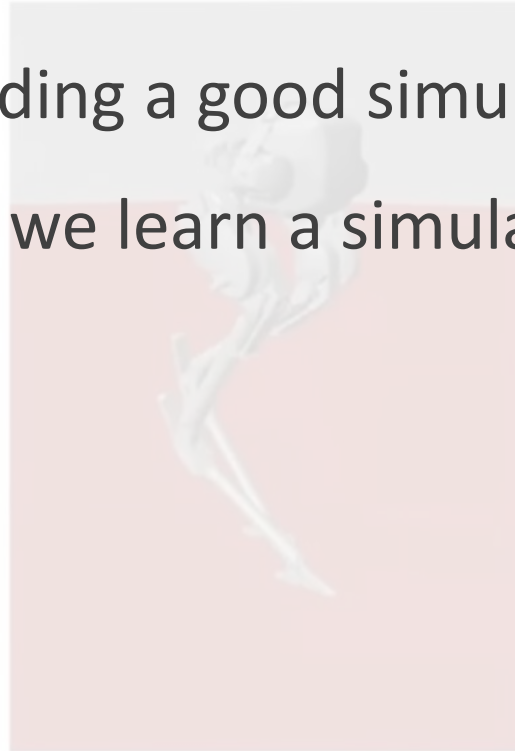

# Sim-to-Real

---





Simulation  
(Low-Fidelity)

Building a good simulator is hard  
Can we learn a simulator?



Simulation  
(High-Fidelity)




Real World

Reinforcement Learning for Robust Parameterized Locomotion Control of Bipedal Robots  
[Li et al. 2021]


# Dynamics Model

---



# Why Learn a Dynamics Model?

---



Simple Dynamics




Complex Dynamics

# Dynamics Model

---

- Learn a dynamics model:


$$f(\mathbf{s}'|\mathbf{s}, \mathbf{a}) \approx p(\mathbf{s}'|\mathbf{s}, \mathbf{a})$$



# Learning Dynamics Model

---

- Collect data with a base policy  $\pi_0$




- Dataset:  $\mathcal{D} = \{(s_i, a_i, s')\}$
- Fit a dynamics model via supervised learning

$$\arg \max_f \mathbb{E}_{(s, a, s') \sim \mathcal{D}} [\log f(s' | s, a)]$$

# Model-Based RL

---

- Collect data with a base policy  $\pi_0$



- Dataset:  $\mathcal{D} = \{(s_i, a_i, s')\}$
- Fit a dynamics model via supervised learning


$$\arg \max_f \mathbb{E}_{(s, a, s') \sim \mathcal{D}} [\log f(s' | s, a)]$$

- Train new policy  $\pi$  by simulating with  $f(s' | s, a)$

# Problem

---


- Reward: climb as high as possible



# Problem

---

- Reward: climb as high as possible






# Problem

---


- Reward: climb as high as possible



# Problem

---


- Reward: climb as high as possible



# Problem

---


- Reward: climb as high as possible



# Problem

---


- Reward: climb as high as possible



# Problem

---


- Reward: climb as high as possible



# Problem

---

- Reward: climb as high as possible



# Distribution Shift


- Data distribution is different from the policy's distribution

$$\mathcal{D} \sim p(\mathbf{s}, \mathbf{a}|\pi_0) \neq p(\mathbf{s}, \mathbf{a}|\pi)$$

- Model  $f(\mathbf{s}'|\mathbf{s}, \mathbf{a})$  trained on  $\mathcal{D}$ 
  - Low error under  $p(\mathbf{s}, \mathbf{a}|\pi_0)$
  - High error under  $p(\mathbf{s}, \mathbf{a}|\pi)$

- Can we make


$$p(\mathbf{s}, \mathbf{a}|\pi_0) = p(\mathbf{s}, \mathbf{a}|\pi) ?$$



# Model-Based RL

---

- Collect data with a base policy  $\pi_0$



- Dataset:  $\mathcal{D} = \{(s_i, a_i, s')\}$
- Fit a dynamics model via supervised learning

$$\arg \max_f \mathbb{E}_{(s,a,s') \sim \mathcal{D}} [\log f(s'|s, a)]$$

- Train new policy  $\pi$  by simulating with  $f$



# DYNA

---

---

## ALGORITHM: DYNA

---

- 1:  $\pi^0 \leftarrow$  initialize policy
  - 2:  $\mathcal{D} \leftarrow \{\emptyset\}$  initialize dataset
  
  - 3: **for** iteration  $k = 0, \dots, n - 1$  **do**
  - 4:   Sample trajectory  $\tau$  according to  $\pi^k(\mathbf{a}|\mathbf{s})$
  - 5:   Add transitions to dataset  $\mathcal{D} = \mathcal{D} \cup \{(\mathbf{s}_i, \mathbf{a}_i, \mathbf{s}'_i)\}$
  
  - 6:   Fit dynamics model:  
     $f = \arg \max_f \mathbb{E}_{(\mathbf{s}, \mathbf{a}, \mathbf{s}') \sim \mathcal{D}} [\log f(\mathbf{s}'|\mathbf{s}, \mathbf{a})]$
  
  - 7:    $\pi^{k+1} \leftarrow$  train policy by simulating rollouts with  $f(\mathbf{s}'|\mathbf{s}, \mathbf{a})$
  - 8: **end for**
  
  - 9: return  $\pi^n$
- 

Dyna, an Integrated Architecture for Learning, Planning, and Reacting  
[Sutton 1991]

# DYNA

---

---

## ALGORITHM: DYNA

---

- 1:  $\pi^0 \leftarrow$  initialize policy
  - 2:  $\mathcal{D} \leftarrow \{\emptyset\}$  initialize dataset
  - 3: **for** iteration  $k = 0, \dots, n - 1$  **do**
  - 4:   Sample trajectory  $\tau$  according to  $\pi^k(\mathbf{a}|\mathbf{s})$
  - 5:   Add transitions to dataset  $\mathcal{D} = \mathcal{D} \cup \{(\mathbf{s}_i, \mathbf{a}_i, \mathbf{s}'_i)\}$
  - 6:   Fit dynamics model:  
     $f = \arg \max_f \mathbb{E}_{(\mathbf{s}, \mathbf{a}, \mathbf{s}') \sim \mathcal{D}} [\log f(\mathbf{s}'|\mathbf{s}, \mathbf{a})]$
  - 7:    $\pi^{k+1} \leftarrow$  train policy by simulating rollouts with  $f(\mathbf{s}'|\mathbf{s}, \mathbf{a})$
  - 8: **end for**
  - 9: **return**  $\pi^n$
- 

Dyna, an Integrated Architecture for Learning, Planning, and Reacting  
[Sutton 1991]

# DYNA

---

---

## ALGORITHM: DYNA

---

- 1:  $\pi^0 \leftarrow$  initialize policy
  - 2:  $\mathcal{D} \leftarrow \{\emptyset\}$  initialize dataset
  - 3: **for** iteration  $k = 0, \dots, n - 1$  **do**
  - 4:   Sample trajectory  $\tau$  according to  $\pi^k(\mathbf{a}|\mathbf{s})$
  - 5:   Add transitions to dataset  $\mathcal{D} = \mathcal{D} \cup \{(\mathbf{s}_i, \mathbf{a}_i, \mathbf{s}'_i)\}$
  - 6:   Fit dynamics model:  
     $f = \arg \max_f \mathbb{E}_{(\mathbf{s}, \mathbf{a}, \mathbf{s}') \sim \mathcal{D}} [\log f(\mathbf{s}'|\mathbf{s}, \mathbf{a})]$
  - 7:    $\pi^{k+1} \leftarrow$  train policy by simulating rollouts with  $f(\mathbf{s}'|\mathbf{s}, \mathbf{a})$
  - 8: **end for**
  - 9: **return**  $\pi^n$
- 

Dyna, an Integrated Architecture for Learning, Planning, and Reacting  
[Sutton 1991]

# DYNA

---

---

## ALGORITHM: DYNA

---

- 1:  $\pi^0 \leftarrow$  initialize policy
  - 2:  $\mathcal{D} \leftarrow \{\emptyset\}$  initialize dataset
  
  - 3: **for** iteration  $k = 0, \dots, n - 1$  **do**
  - 4:   Sample trajectory  $\tau$  according to  $\pi^k(\mathbf{a}|\mathbf{s})$
  - 5:   Add transitions to dataset  $\mathcal{D} = \mathcal{D} \cup \{(\mathbf{s}_i, \mathbf{a}_i, \mathbf{s}'_i)\}$
  
  - 6:   Fit dynamics model:  
     $f = \arg \max_f \mathbb{E}_{(\mathbf{s}, \mathbf{a}, \mathbf{s}') \sim \mathcal{D}} [\log f(\mathbf{s}'|\mathbf{s}, \mathbf{a})]$
  
  - 7:    $\pi^{k+1} \leftarrow$  train policy by simulating rollouts with  $f(\mathbf{s}'|\mathbf{s}, \mathbf{a})$
  - 8: **end for**
  
  - 9: return  $\pi^n$
- 

Dyna, an Integrated Architecture for Learning, Planning, and Reacting  
[Sutton 1991]

# DYNA

---

---

## ALGORITHM: DYNA

---

- 1:  $\pi^0 \leftarrow$  initialize policy
  - 2:  $\mathcal{D} \leftarrow \{\emptyset\}$  initialize dataset
  
  - 3: **for** iteration  $k = 0, \dots, n - 1$  **do**
  - 4:   Sample trajectory  $\tau$  according to  $\pi^k(\mathbf{a}|\mathbf{s})$
  - 5:   Add transitions to dataset  $\mathcal{D} = \mathcal{D} \cup \{(\mathbf{s}_i, \mathbf{a}_i, \mathbf{s}'_i)\}$
  - 6:   Fit dynamics model:  
     $f = \arg \max_f \mathbb{E}_{(\mathbf{s}, \mathbf{a}, \mathbf{s}') \sim \mathcal{D}} [\log f(\mathbf{s}'|\mathbf{s}, \mathbf{a})]$
  - 7:    $\pi^{k+1} \leftarrow$  train policy by simulating rollouts with  $f(\mathbf{s}'|\mathbf{s}, \mathbf{a})$
  - 8: **end for**
  
  - 9: return  $\pi^n$
- 

Dyna, an Integrated Architecture for Learning, Planning, and Reacting  
[Sutton 1991]

# DYNA

---

---

## ALGORITHM: DYNA

---

- 1:  $\pi^0 \leftarrow$  initialize policy
  - 2:  $\mathcal{D} \leftarrow \{\emptyset\}$  initialize dataset
  
  - 3: **for** iteration  $k = 0, \dots, n - 1$  **do**
  - 4:   Sample trajectory  $\tau$  according to  $\pi^k(\mathbf{a}|\mathbf{s})$
  - 5:   Add transitions to dataset  $\mathcal{D} = \mathcal{D} \cup \{(\mathbf{s}_i, \mathbf{a}_i, \mathbf{s}'_i)\}$
  
  - 6:   Fit dynamics model:  
     $f = \arg \max_f \mathbb{E}_{(\mathbf{s}, \mathbf{a}, \mathbf{s}') \sim \mathcal{D}} [\log f(\mathbf{s}'|\mathbf{s}, \mathbf{a})]$
  - 7:    $\pi^{k+1} \leftarrow$  train policy by simulating rollouts with  $f(\mathbf{s}'|\mathbf{s}, \mathbf{a})$
  - 8: **end for**
  
  - 9: return  $\pi^n$
- 

Dyna, an Integrated Architecture for Learning, Planning, and Reacting  
[Sutton 1991]

# DYNA

---

---

## ALGORITHM: DYNA

---

- 1:  $\pi^0 \leftarrow$  initialize policy
  - 2:  $\mathcal{D} \leftarrow \{\emptyset\}$  initialize dataset
  
  - 3: **for** iteration  $k = 0, \dots, n - 1$  **do**
  - 4:   Sample trajectory  $\tau$  according to  $\pi^k(\mathbf{a}|\mathbf{s})$
  - 5:   Add transitions to dataset  $\mathcal{D} = \mathcal{D} \cup \{(\mathbf{s}_i, \mathbf{a}_i, \mathbf{s}'_i)\}$
  
  - 6:   Fit dynamics model:  
       $f = \arg \max_f \mathbb{E}_{(\mathbf{s}, \mathbf{a}, \mathbf{s}') \sim \mathcal{D}} [\log f(\mathbf{s}'|\mathbf{s}, \mathbf{a})]$
  
  - 7:    $\pi^{k+1} \leftarrow$  train policy by simulating rollouts with  $f(\mathbf{s}'|\mathbf{s}, \mathbf{a})$
  - 8: **end for**
  
  - 9: return  $\pi^n$
- 

Dyna, an Integrated Architecture for Learning, Planning, and Reacting  
[Sutton 1991]

# DYNA

---

---

## ALGORITHM: DYNA

---

- 1:  $\pi^0 \leftarrow$  initialize policy
  - 2:  $\mathcal{D} \leftarrow \{\emptyset\}$  initialize dataset
  
  - 3: **for** iteration  $k = 0, \dots, n - 1$  **do**
  - 4:   Sample trajectory  $\tau$  according to  $\pi^k(\mathbf{a}|\mathbf{s})$
  - 5:   Add transitions to dataset  $\mathcal{D} = \mathcal{D} \cup \{(\mathbf{s}_i, \mathbf{a}_i, \mathbf{s}'_i)\}$
  
  - 6:   Fit dynamics model:  
       $f = \arg \max_f \mathbb{E}_{(\mathbf{s}, \mathbf{a}, \mathbf{s}') \sim \mathcal{D}} [\log f(\mathbf{s}'|\mathbf{s}, \mathbf{a})]$
  
  - 7:    $\pi^{k+1} \leftarrow$  train policy by simulating rollouts with  $f(\mathbf{s}'|\mathbf{s}, \mathbf{a})$
  - 8: **end for**
  
  - 9: return  $\pi^n$
- 

Dyna, an Integrated Architecture for Learning, Planning, and Reacting  
[Sutton 1991]



# DYNA

---

## ALGORITHM: DYNA

---

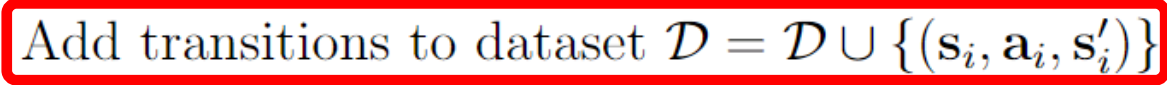
- 1:  $\pi^0 \leftarrow$  initialize policy
  - 2:  $\mathcal{D} \leftarrow \{\emptyset\}$  initialize dataset
  
  - 3: **for** iteration  $k = 0, \dots, n - 1$  **do**
  - 4:   Sample trajectory  $\tau$  according to  $\pi^k(\mathbf{a}|\mathbf{s})$
  - 5:   Add transitions to dataset  $\mathcal{D} = \mathcal{D} \cup \{(\mathbf{s}_i, \mathbf{a}_i, \mathbf{s}'_i)\}$
  
  - 6:   Fit dynamics model:  
     $f = \arg \max_f \mathbb{E}_{(\mathbf{s}, \mathbf{a}, \mathbf{s}') \sim \mathcal{D}} [\log f(\mathbf{s}'|\mathbf{s}, \mathbf{a})]$
  - 7:    $\pi^{k+1} \leftarrow$  train policy by simulating rollouts with  $f(\mathbf{s}'|\mathbf{s}, \mathbf{a})$
  - 8: **end for**
  
  - 9: return  $\pi^n$
- (e.g. policy gradient, Q-learning, SAC, etc.)
-

# DYNA

---

## ALGORITHM: DYNA



---

- 1:  $\pi^0 \leftarrow$  initialize policy
  - 2:  $\mathcal{D} \leftarrow \{\emptyset\}$  initialize dataset
  - 3: **for** iteration  $k = 0, \dots, n - 1$  **do**
  - 4:   Sample trajectory  $\tau$  according to  $\pi^k(\mathbf{a}|\mathbf{s})$
  - 5:   Add transitions to dataset  $\mathcal{D} = \mathcal{D} \cup \{(\mathbf{s}_i, \mathbf{a}_i, \mathbf{s}'_i)\}$  
  - 6:   Fit dynamics model:  
     $f = \arg \max_f \mathbb{E}_{(\mathbf{s}, \mathbf{a}, \mathbf{s}') \sim \mathcal{D}} [\log f(\mathbf{s}'|\mathbf{s}, \mathbf{a})]$
  - 7:    $\pi^{k+1} \leftarrow$  train policy by simulating rollouts with  $f(\mathbf{s}'|\mathbf{s}, \mathbf{a})$
  - 8: **end for**
  - 9: **return**  $\pi^n$
- 

keep data from  
all iterations

# Model Representation

- How do we represent  $f(s'|s, a)$ ?
- MDP with small discrete states and actions  $\rightarrow$  lookup table




# Deterministic Models

---

- How do we represent  $f(s'|s, a)$ ?

$$\arg \min_f \mathbb{E}_{(s,a,s') \sim \mathcal{D}} [||s' - f(s, a)||^2]$$

What if the dynamics  
are stochastic?




# Stochastic Models

---

- How do we represent  $f(s'|s, a)$ ?

$$\arg \max_f \mathbb{E}_{(s,a,s') \sim \mathcal{D}} [\log f(s'|s, a)]$$



# Stochastic Models

---

- How do we represent  $f(s'|s, a)$ ?

$$\arg \max_f \mathbb{E}_{(s,a,s') \sim \mathcal{D}} [\underbrace{\log f(s'|s, a)}]$$

## Conditional Generative Model

- Variational Autoencoders (VAEs)
- Generative Adversarial Networks (GANs)
- Flow Models
- Diffusion Models
- Etc.

# Reward Model

---

- If reward function is unknown, augment model to predict both states and rewards
- For most tasks, reward function is available/specified by a human

dynamics model

$$f(\mathbf{s}'|\mathbf{s}, \mathbf{a})$$

reward model

$$h(r|\mathbf{s}, \mathbf{a}, \mathbf{s}')$$

# Model-Based Rollout

---

---

## ALGORITHM: DYNA

---



- 1:  $\pi^0 \leftarrow$  initialize policy
  - 2:  $\mathcal{D} \leftarrow \{\emptyset\}$  initialize dataset
  - 3: **for** iteration  $k = 0, \dots, n - 1$  **do**
  - 4:   Sample trajectory  $\tau$  according to  $\pi^k(\mathbf{a}|\mathbf{s})$
  - 5:   Add transitions to dataset  $\mathcal{D} = \mathcal{D} \cup \{(\mathbf{s}_i, \mathbf{a}_i, \mathbf{s}'_i)\}$
  - 6:   Fit dynamics model:  
     $f = \arg \max_f \mathbb{E}_{(\mathbf{s}, \mathbf{a}, \mathbf{s}') \sim \mathcal{D}} [\log f(\mathbf{s}'|\mathbf{s}, \mathbf{a})]$
  - 7:    $\pi^{k+1} \leftarrow$  train policy by simulating rollouts with  $f(\mathbf{s}'|\mathbf{s}, \mathbf{a})$
  - 8: **end for**
  - 9: **return**  $\pi^n$
- 

generate trajectories  
with model



# Model-Based Rollout


---



# Drift


---

- Same action sequence in the real env and the model can lead to very different trajectories
- Autoregressive model  $\rightarrow$  compounding error




# Drift


---




# Model-Based Rollout



# Model-Based Rollout



Generate shorter rollouts




# Model-Based Rollout

---

— — — Real  
- - - Model

Generate shorter rollouts  
from different real states




# Model-Based RL




(a) Example task domains.

model-based



(b) NAF and DDPG on multi-target reacher.

model-based



(c) NAF and DDPG on peg insertion.

Continuous Deep Q-Learning with Model-based Acceleration  
[Gu et al. 2016]

# DYNA

---

## ALGORITHM: DYNA

---

- 1:  $\pi^0 \leftarrow$  initialize policy
  - 2:  $\mathcal{D} \leftarrow \{\emptyset\}$  initialize dataset
  
  - 3: **for** iteration  $k = 0, \dots, n - 1$  **do**
  - 4:   Sample trajectory  $\tau$  according to  $\pi^k(\mathbf{a}|\mathbf{s})$
  - 5:   Add transitions to dataset  $\mathcal{D} = \mathcal{D} \cup \{(\mathbf{s}_i, \mathbf{a}_i, \mathbf{s}'_i)\}$
  
  - 6:   Fit dynamics model:  
     $f = \arg \max_f \mathbb{E}_{(\mathbf{s}, \mathbf{a}, \mathbf{s}') \sim \mathcal{D}} [\log f(\mathbf{s}'|\mathbf{s}, \mathbf{a})]$
  - 7:    $\pi^{k+1} \leftarrow$  train policy by simulating rollouts with  $f(\mathbf{s}'|\mathbf{s}, \mathbf{a})$
  - 8: **end for**
  
  - 9: return  $\pi^n$
- use any RL algorithm  
(e.g. policy gradient, Q-learning, SAC, etc.)
- 


Dyna, an Integrated Architecture for Learning, Planning, and Reacting  
[Sutton 1991]



# Differentiable Dynamics

---


$$\arg \max_{\pi} \mathbb{E}_{\tau \sim \underline{p(\tau|\pi)}} \left[ \sum_{t=0}^{T-1} \gamma^t r_t \right]$$



$$\arg \max_{\pi} \mathbb{E}_{\tau \sim \underline{f(\tau|\pi)}} \left[ \sum_{t=0}^{T-1} \gamma^t r_t \right]$$

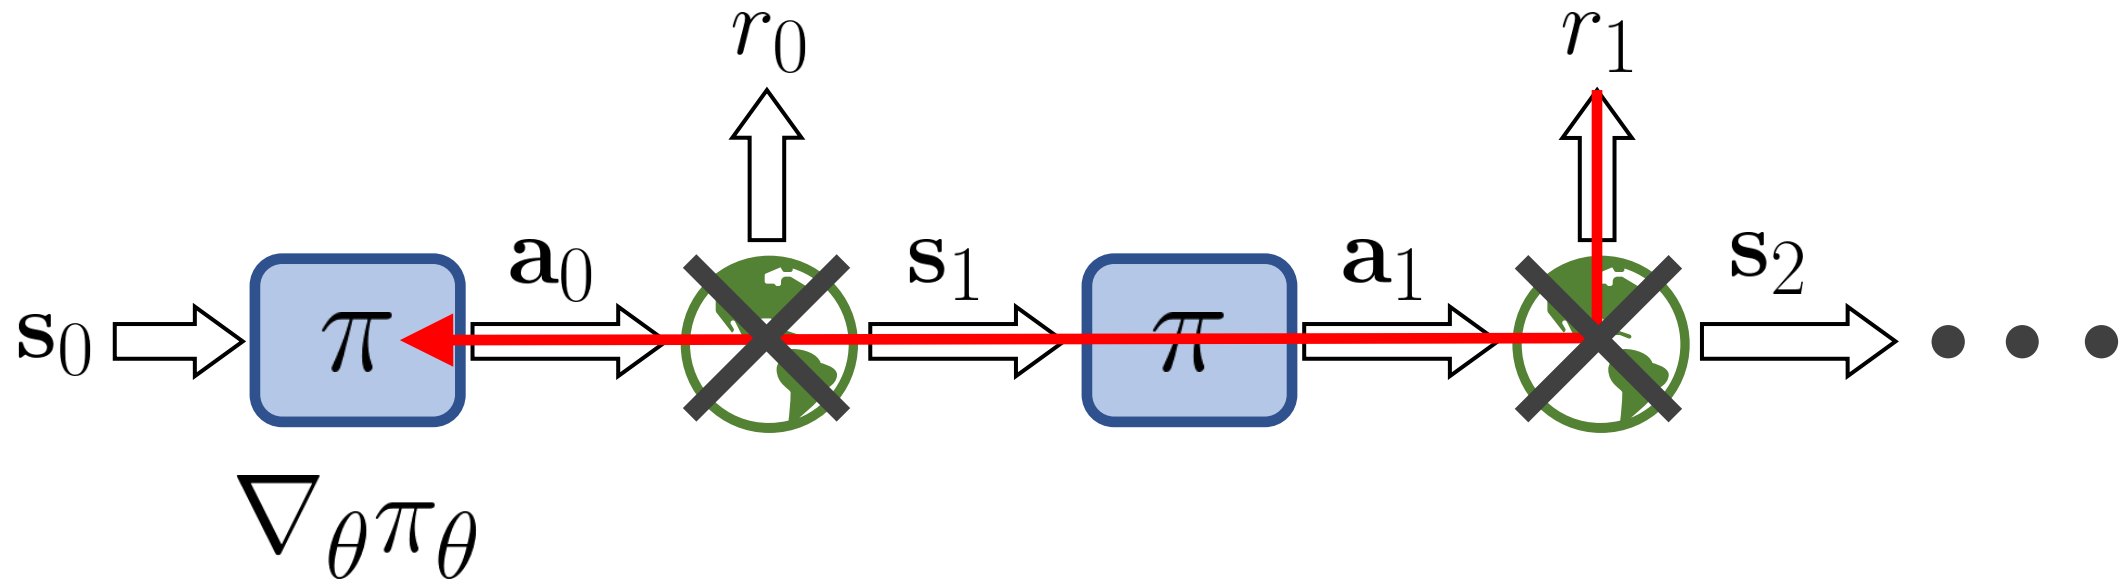
use any RL algorithm  
(e.g. policy gradient, Q-learning, SAC, etc.)

# Differentiable Dynamics




$$\frac{\partial r_1}{\partial \theta} = \overset{\text{red X}}{\frac{\partial r_1}{\partial \mathbf{a}_1}} \overset{\text{green check}}{\frac{\partial \mathbf{a}_1}{\partial \theta}} + \overset{\text{red X}}{\frac{\partial r_1}{\partial \mathbf{a}_1}} \overset{\text{green check}}{\frac{\partial \mathbf{a}_1}{\partial s_1}} \overset{\text{red X}}{\frac{\partial s_1}{\partial \mathbf{a}_0}} \overset{\text{green check}}{\frac{\partial \mathbf{a}_0}{\partial \theta}} + \dots$$

# Differentiable Dynamics



$$\frac{\partial r_1}{\partial \theta} = \overset{\checkmark}{\frac{\partial r_1}{\partial \mathbf{a}_1}} \overset{\checkmark}{\frac{\partial \mathbf{a}_1}{\partial \theta}} + \overset{\checkmark}{\frac{\partial r_1}{\partial \mathbf{a}_1}} \overset{\checkmark}{\frac{\partial \mathbf{a}_1}{\partial s_1}} \overset{\times}{\frac{\partial s_1}{\partial \mathbf{a}_0}} \overset{\checkmark}{\frac{\partial \mathbf{a}_0}{\partial \theta}} + \dots$$


# Differentiable Dynamics



$$\frac{\partial r_1}{\partial \theta} = \frac{\partial r_1}{\partial \mathbf{a}_1} \frac{\partial \mathbf{a}_1}{\partial \theta} + \frac{\partial r_1}{\partial \mathbf{a}_1} \frac{\partial \mathbf{a}_1}{\partial s_1} \boxed{\frac{\partial s_1}{\partial \mathbf{a}_0}} \frac{\partial \mathbf{a}_0}{\partial \theta} + \dots$$

Dynamics

# Differentiable Dynamics




Fully Differentiable!

$$\frac{\partial r_1}{\partial \theta} = \frac{\partial r_1}{\partial \mathbf{a}_1} \frac{\partial \mathbf{a}_1}{\partial \theta} + \frac{\partial r_1}{\partial \mathbf{a}_1} \frac{\partial \mathbf{a}_1}{\partial s_1} \frac{\partial s_1}{\partial \mathbf{a}_0} \frac{\partial \mathbf{a}_0}{\partial \theta} + \dots$$

# Differentiable Dynamics


---

$$\arg \max_{\pi} \mathbb{E}_{\tau \sim f(\tau|\pi)} \left[ \sum_{t=0}^{T-1} \gamma^t r_t \right]$$


Compute gradients using autodiff  
and solve with gradient ascent

# Differentiable Dynamics

---




SuperTrack: Motion Tracking for Physically Simulated Characters Using Supervised Learning  
[Fussell et al. 2021]


# Differentiable Dynamics



A1 Quadruped  
Walking



UR5 Multi-Object  
Visual Pick Place



XArm Visual Pick  
and Place



Sphero Ollie Visual  
Navigation

DayDreamer: World Models for Physical Robot Learning  
[Wu et al. 2022]



# Differentiable Dynamics

---



DayDreamer: World Models for Physical Robot Learning  
[Wu et al. 2022]

# Differentiable Dynamics

---



DayDreamer: World Models for Physical Robot Learning  
[Wu et al. 2022]

# Model Exploitation

---

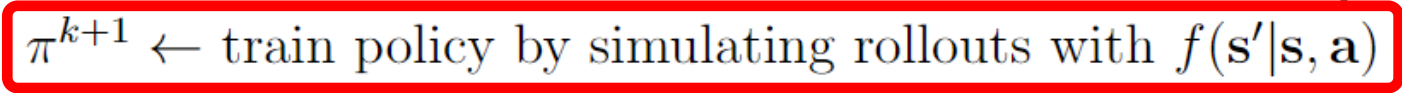
---

## ALGORITHM: DYNA

---

- 1:  $\pi^0 \leftarrow$  initialize policy
  - 2:  $\mathcal{D} \leftarrow \{\emptyset\}$  initialize dataset
  
  - 3: **for** iteration  $k = 0, \dots, n - 1$  **do**
  - 4:   Sample trajectory  $\tau$  according to  $\pi^k(\mathbf{a}|\mathbf{s})$
  - 5:   Add transitions to dataset  $\mathcal{D} = \mathcal{D} \cup \{(\mathbf{s}_i, \mathbf{a}_i, \mathbf{s}'_i)\}$
  
  - 6:   Fit dynamics model:  
       $f = \arg \max_f \mathbb{E}_{(\mathbf{s}, \mathbf{a}, \mathbf{s}') \sim \mathcal{D}} [\log f(\mathbf{s}'|\mathbf{s}, \mathbf{a})]$
  
  - 7:    $\pi^{k+1} \leftarrow$  train policy by simulating rollouts with  $f(\mathbf{s}'|\mathbf{s}, \mathbf{a})$
  - 8: **end for**
  
  - 9: **return**  $\pi^n$
- 


Policy can exploit  
errors in model




# 2 Types of Uncertainty

---

Aleatoric  
(Statistical Uncertainty)




Epistemic  
(Model Uncertainty)



# 2 Types of Uncertainty

---


Aleatoric  
(Statistical Uncertainty)




# 2 Types of Uncertainty

---


Aleatoric  
(Statistical Uncertainty)




Epistemic  
(Model Uncertainty)



# 2 Types of Uncertainty




Epistemic  
(Model Uncertainty)




no data

# 2 Types of Uncertainty

Aleatoric  
(Statistical Uncertainty)



Epistemic  
(Model Uncertainty)



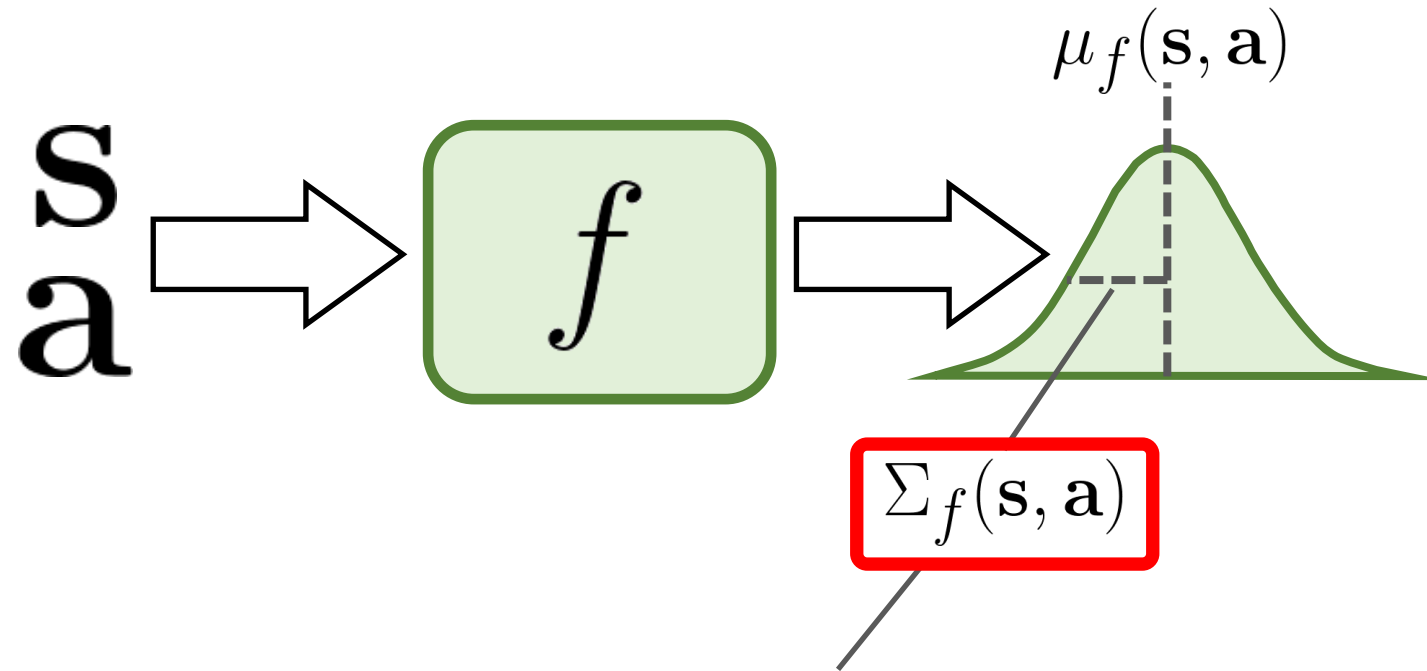
Policy can exploit model uncertainty



# Uncertainty Estimation

---

- Can we estimate the model uncertainty?



This only estimates  
*statistical* uncertainty


# Uncertainty Estimation

---

- Can we estimate the model uncertainty?
- Bayesian inference:

$$\underline{p(f|\mathcal{D})}$$

What is the likelihood of  
a function given the data




# Uncertainty Estimation

---

- Can we estimate the model uncertainty?
- Bayesian inference:

$$\underline{p(f|\mathcal{D})}$$

What is the likelihood of  
a function given the data



High Likelihood


# Uncertainty Estimation

---


- Can we estimate the model uncertainty?
- Bayesian inference:

$$\underline{p(f|\mathcal{D})}$$

What is the likelihood of  
a function given the data



High Likelihood



Low Likelihood

# Uncertainty Estimation

---

- Can we estimate the model uncertainty?
- Bayesian inference:

$$p(f|\mathcal{D}) = \frac{p(f, \mathcal{D})}{p(\mathcal{D})}$$

# Uncertainty Estimation

---

- Can we estimate the model uncertainty?
- Bayesian inference:

$$\begin{aligned} p(f|\mathcal{D}) &= \frac{p(f, \mathcal{D})}{p(\mathcal{D})} \\ &= \frac{p(\mathcal{D}|f)p(f)}{p(\mathcal{D})} \end{aligned}$$

# Supervised Learning

---

$$\arg \max_f \log p(f|\mathcal{D}) = \arg \max_f \log \frac{p(\mathcal{D}|f)p(f)}{p(\mathcal{D})}$$

# Supervised Learning

---


$$\begin{aligned} \arg \max_f \log \underbrace{p(f|\mathcal{D})}_{\text{Posterior}} &= \arg \max_f \log \frac{p(\mathcal{D}|f)p(f)}{p(\mathcal{D})} \\ &= \arg \max_f \underbrace{\log p(\mathcal{D}|f)}_{\text{Likelihood}} + \underbrace{\log p(f)}_{\text{Prior}} - \underbrace{\log p(\mathcal{D})}_{\text{Constant}} \end{aligned}$$



# Supervised Learning

---

$$\begin{aligned}\arg \max_f \log p(f|\mathcal{D}) &= \arg \max_f \log \frac{p(\mathcal{D}|f)p(f)}{p(\mathcal{D})} \\ &= \arg \max_f \underbrace{\log p(\mathcal{D}|f)}_{\text{Likelihood}} + \underbrace{\log p(f)} - \underbrace{\log p(\mathcal{D})}\end{aligned}$$




Maximum Likelihood


$$\arg \max_f \mathbb{E}_{(\mathbf{s}, \mathbf{a}, \mathbf{s}') \sim \mathcal{D}} [\log f(\mathbf{s}'|\mathbf{s}, \mathbf{a})]$$

# Supervised Learning

---



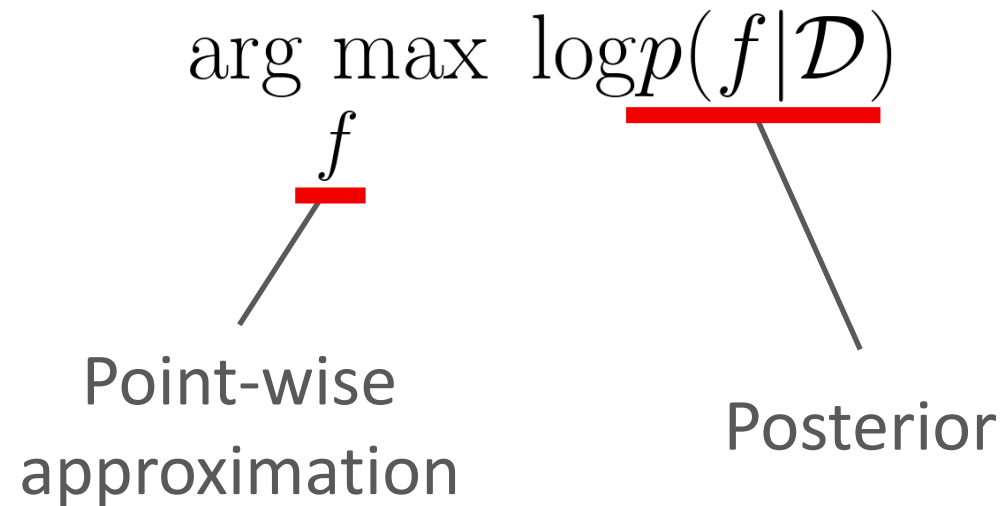
# Supervised Learning



# Uncertainty Estimation


---

- Maximum likelihood only gives a *point-wise* approximation of the posterior
- To estimate model uncertainty, need to approximate the *full* posterior



# Ensemble


- Approximate posterior with ensemble



# Ensemble

---


- Approximate posterior with ensemble
- Models should be consistent under the data distribution



# Ensemble

---

- Approximate posterior with ensemble
- Models should be consistent under the data distribution
- Models will hopefully disagree on out-of-distribution samples




# How to train ensemble?

---

## Bootstrapping

- Split dataset into subsets
- Train a separate model for each subset

**×** Reduces data available to train each model






# How to train ensemble?

---

## Bootstrapping

- Split dataset into subsets
- Train a separate model for each subset

 Reduces data available to train each model


## In practice:

- Initialize models with different random parameters
- Train all models using the same dataset
- Stochasticity from SGD leads to diverse models

# How to use ensemble?

---


- Sample random model for every transition



# How to use ensemble?

---


- Sample random model for every transition



# How to use ensemble?

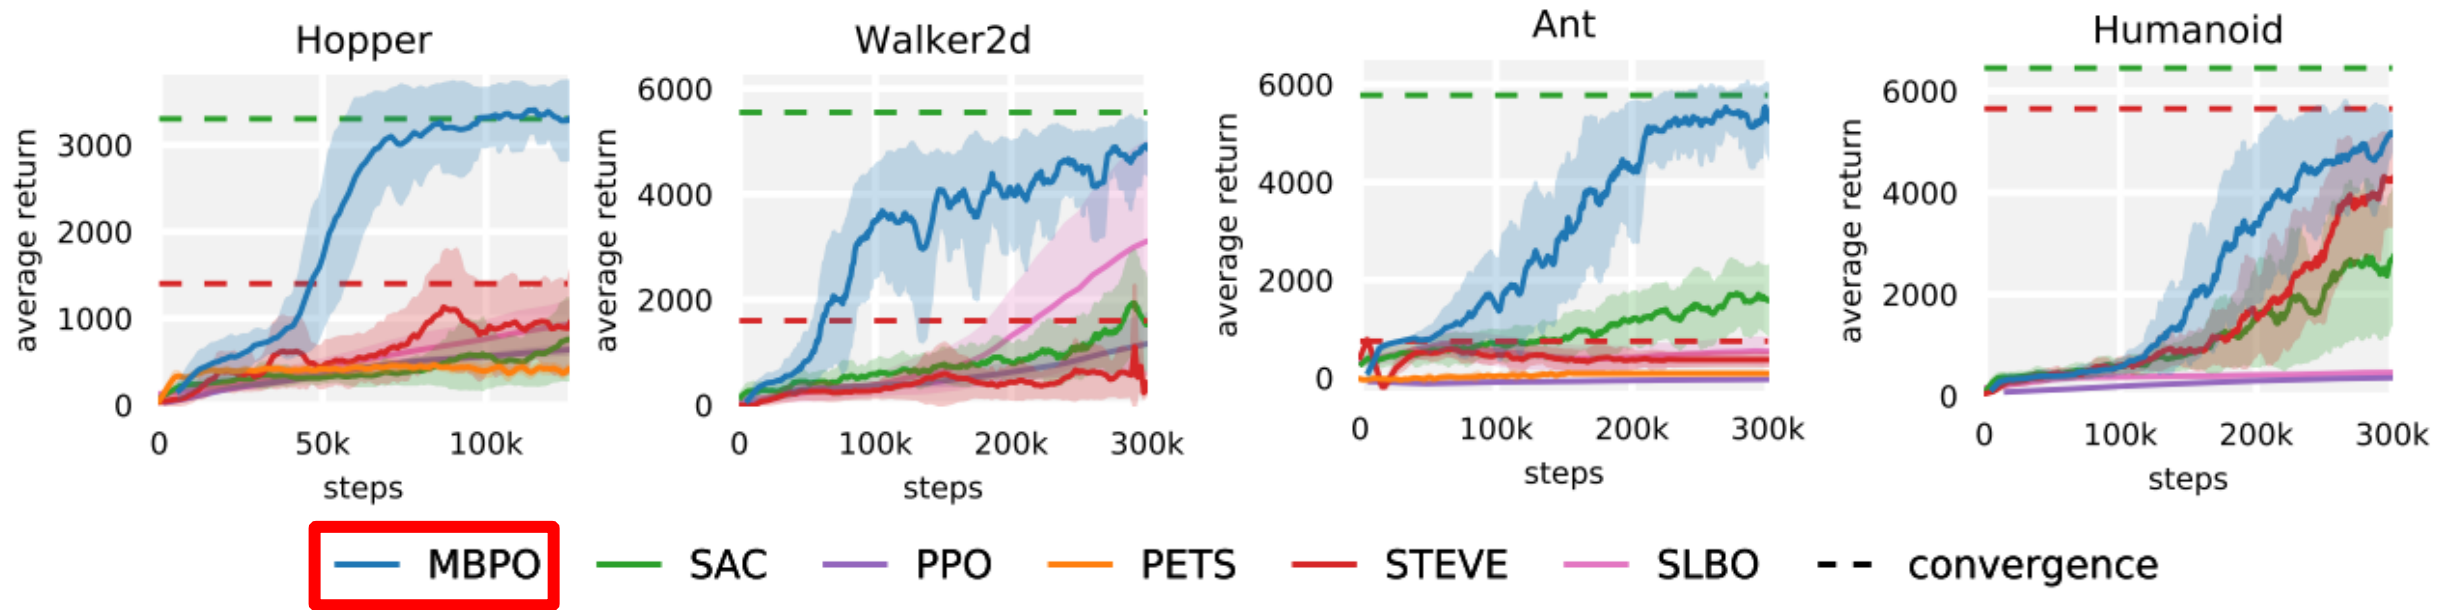
---

- Sample random model for every transition



# How to use ensemble?

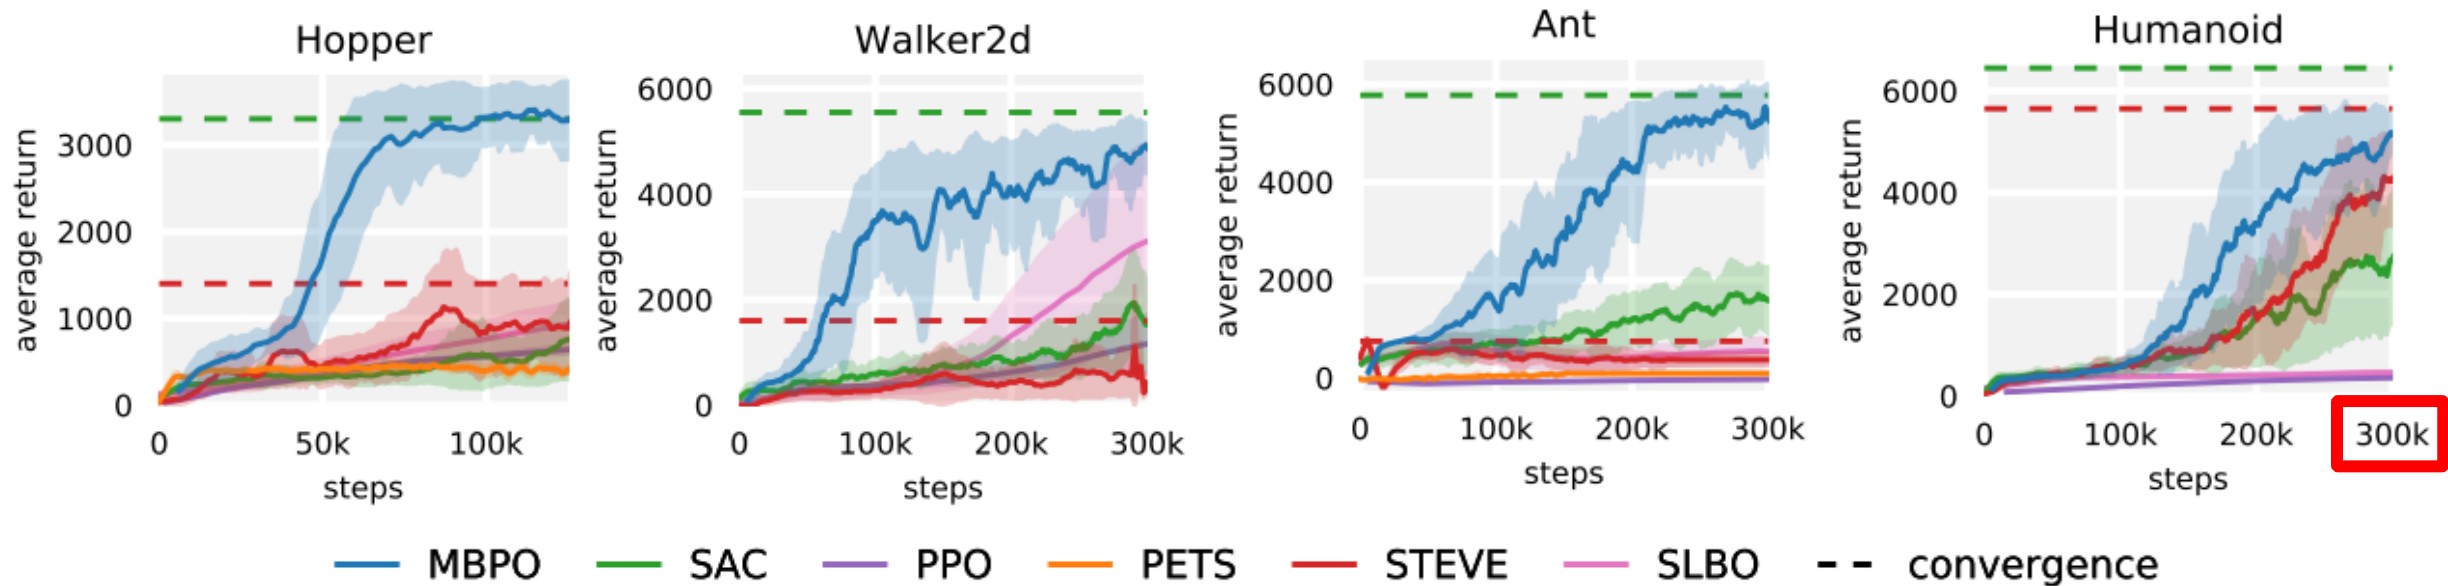
- Sample random model for every transition



When to Trust Your Model: Model-Based Policy Optimization  
[Janner et al. 2019]

# How to use ensemble?

- Sample random model for every transition



When to Trust Your Model: Model-Based Policy Optimization  
[Janner et al. 2019]

# How to use ensemble?

---

- Sample random model for every transition
- Penalize policy for model disagreement

$$r(\mathbf{s}, \mathbf{a}, \mathbf{s}')$$

# How to use ensemble?

---

- Sample random model for every transition
- Penalize policy for model disagreement

$$r_p(\mathbf{s}, \mathbf{a}, \mathbf{s}') = \begin{cases} -\kappa & \text{if } \underline{d(\mathbf{s}, \mathbf{a})} > \alpha \\ r(\mathbf{s}, \mathbf{a}, \mathbf{s}') & \text{otherwise} \end{cases}$$



# How to use ensemble?

---

- Sample random model for every transition
- Penalize policy for model disagreement

$$r_p(\mathbf{s}, \mathbf{a}, \mathbf{s}') = \begin{cases} -\kappa & \text{if } d(\mathbf{s}, \mathbf{a}) > \underline{\alpha} \\ r(\mathbf{s}, \mathbf{a}, \mathbf{s}') & \text{otherwise} \end{cases}$$

# How to use ensemble?

---

- Sample random model for every transition
- Penalize policy for model disagreement

$$r_p(\mathbf{s}, \mathbf{a}, \mathbf{s}') = \begin{cases} -\kappa & \text{if } d(\mathbf{s}, \mathbf{a}) > \alpha \\ r(\mathbf{s}, \mathbf{a}, \mathbf{s}') & \text{otherwise} \end{cases}$$

# How to use ensemble?

---

- Sample random model for every transition
- Penalize policy for model disagreement

$$r_p(\mathbf{s}, \mathbf{a}, \mathbf{s}') = \begin{cases} -\kappa & \text{if } d(\mathbf{s}, \mathbf{a}) > \alpha \\ r(\mathbf{s}, \mathbf{a}, \mathbf{s}') & \text{otherwise} \end{cases}$$


Model disagreement:

$$d(\mathbf{s}, \mathbf{a}) = \max_{i,j} D(f_i(\cdot|\mathbf{s}, \mathbf{a}), f_j(\cdot|\mathbf{s}, \mathbf{a}))$$

# How to use ensemble?

---

- Sample random model for every transition
- Penalize policy for model disagreement
- Termination based on disagreement



# Uncertainty Estimation

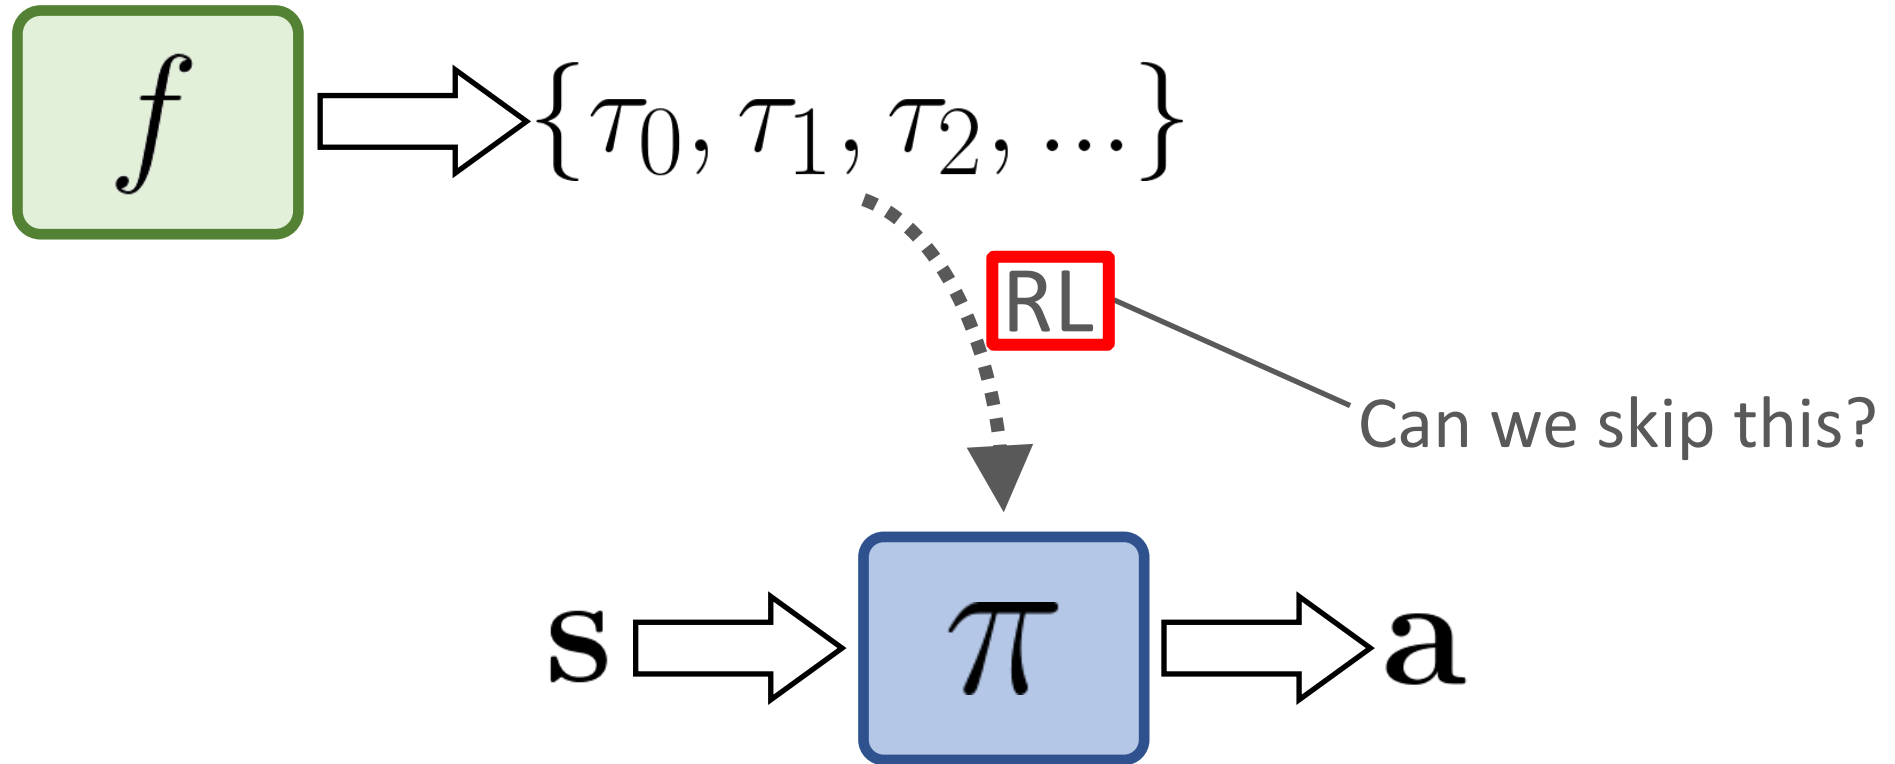
---

- Ensembles
- Bayesian Neural Networks
- Dropout
- Normalized Maximum Likelihood
- Test Time Augmentation
- Etc...

# Model-Predictive Control


# Model-Based Policy Learning

---



# Q-Learning


---






# Online Planning

---




# Online Planning

---




# Online Planning

---



# Online Planning

---




# Online Planning

---

- Use dynamics model to predict expected return of every action

$$\arg \max_{\mathbf{a}_{0:k}} \mathbb{E}_{\tau \sim f(\tau | \mathbf{s}_0 = \mathbf{s}, \mathbf{a}_{0:k})} [R(\tau)]$$




# Online Planning

---

- Use dynamics model to predict expected return of every action


$$\arg \max_{\mathbf{a}_{0:k}} \mathbb{E}_{\tau \sim f(\tau | \mathbf{s}_0 = \mathbf{s}, \mathbf{a}_{0:k})} [R(\tau)]$$

- Apply optimal action sequence  $\mathbf{a}_{0:k}^*$  in real environment




# Drift

---



# Drift


---






# Drift

---




# Drift

---




# Drift

---




# Drift

---



# Drift

---




# MPC

---

- Use dynamics model to predict expected return of every action


$$\arg \max_{\mathbf{a}_{0:k}} \mathbb{E}_{\tau \sim f(\tau | \mathbf{s}_0 = \mathbf{s}, \mathbf{a}_{0:k})} [R(\tau)]$$

- ~~Apply optimal action sequence  $\mathbf{a}_{0:k}^*$  in real environment~~
- Model Predictive Control (MPC)
  - Apply only the first action in the real environment
  - Replan every timestep




# Drift

---



# Drift


---






# Drift

---




# Drift

---




# Drift

---




# Drift

---



# Drift

---



(Closed-Loop Control)

# MPC

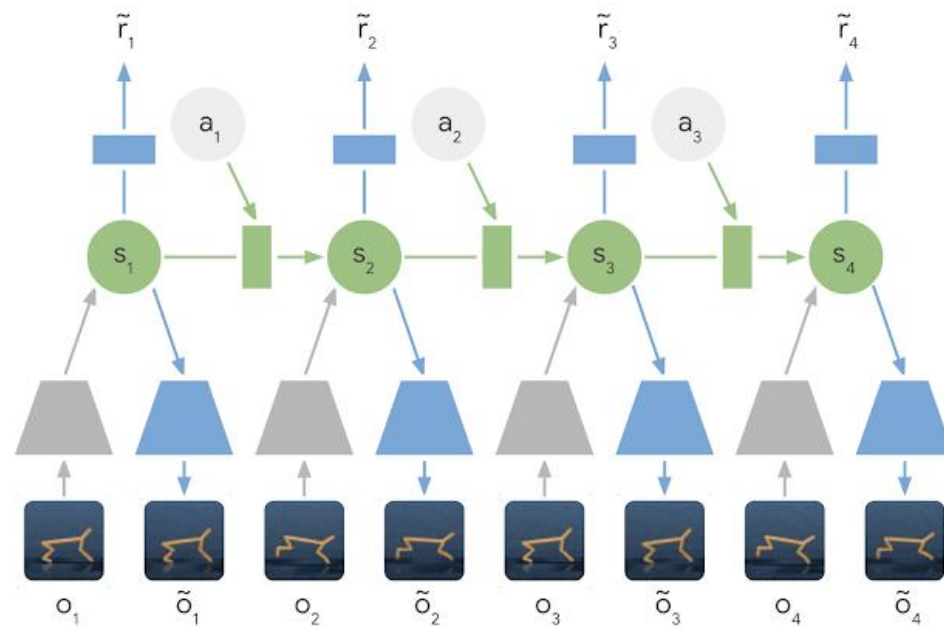

---

- How to solve optimization problem every timestep?

$$\arg \max_{\mathbf{a}_{0:k}} \mathbb{E}_{\tau \sim f(\tau | \mathbf{s}_0 = \mathbf{s}, \mathbf{a}_{0:k})} [R(\tau)]$$

- Black-Box Optimization
  - CEM, random shooting, etc.
- If differentiable model and reward function, use gradient ascent
- Can incorporate other model-based RL improvements
  - Uncertainty estimation, ensembles, etc.


# MPC



Learning Latent Dynamics for Planning from Pixels  
[Hafner et al. 2019]

# MPC



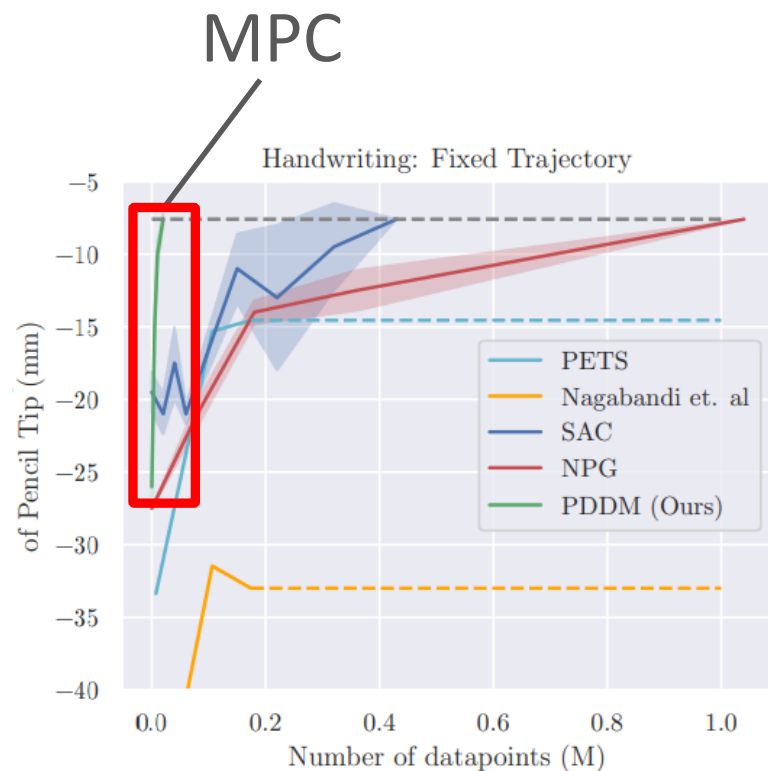
---



Deep Dynamics Models for Learning Dexterous Manipulation  
[Nagabandi et al. 2019]



# MPC



Deep Dynamics Models for Learning Dexterous Manipulation  
[Nagabandi et al. 2019]

# Model-Based RL

---

## **Policy Learning**

- Learn model + policy
- Runtime policy inference is fast
- Policy is task-specific
- Typically better asymptotic performance

## **Online Planning**

- Learn model
- Runtime planning can be slow
- Model can be task-agnostic
- May need many samples during online planning to find good plans

# Summary

---

- Model-Based RL
- DYNA
- Model Representations
- Uncertainty Estimation
- MPC